

Report

Steps:

```
$ unzip 2018114016.zip
```

```
$ cd 2018114016
```

```
$ python3 language_model.py <n-1,2,3> <k/n> corpus.txt
```

A prompt will occur:

```
$ input your sentence: <input test sentence here>
```

Cleaning the data:

Corpus data is cleaned/preprocessed. Numbers are replaced by empty string.

Multiple `\n` are replaced by single `\n`. Parentheses, other brackets are removed.

Alternate sentence endings such as `! , ?` are replaced by `.\n`. Multiple spaces are replaced by single space. This data is then used as corpus and fed to the `grams` function which fills the unigram, bigram and trigram matrices accordingly.

Smoothing:

The idea behind smoothing is to compensate/manage grams which haven't been seen before so the overall probability does not become zero just because a new token occurred.

How kneserney does this is by using something what's called interpolation by absolute discounting. Basically, the Lower order model becomes more necessary/significant only when the count in higher order model and so weights are assigned accordingly. For counts greater than 1, a fixed discount value is subtracted and weights of different order models are assigned as per the count.

How Witten Bell tackles this issue is by doing what is called a backoff. If a higher model has 0 count, we backoff to a lower order model. It tries to capture the idea that if a context appears more number of times, a new token is more likely to appear in that context.

Results

n=2

Sentence	Witten Bell	KneserNey
I am a man	0.000144241470552597	1.68586977709641E-05
The king is not pretty	4.94549823315507E-06	2.66751547008925E-07
Apple is red	0.00748713865752082	Float division error

n=3

Sentence	Witten Bell	KneserNey
I am a man	0.0205246100767598	0.00146986771190593
The king is not pretty	0.00294012874306326	5.58177612116176E-05
Apple is red	0.143269230769231	0.0384615384615385